YuE2: SongEval validation after GRPO

A

CMI-Pref · 270 train / 30 fixed validation conditions · endpoint step 270

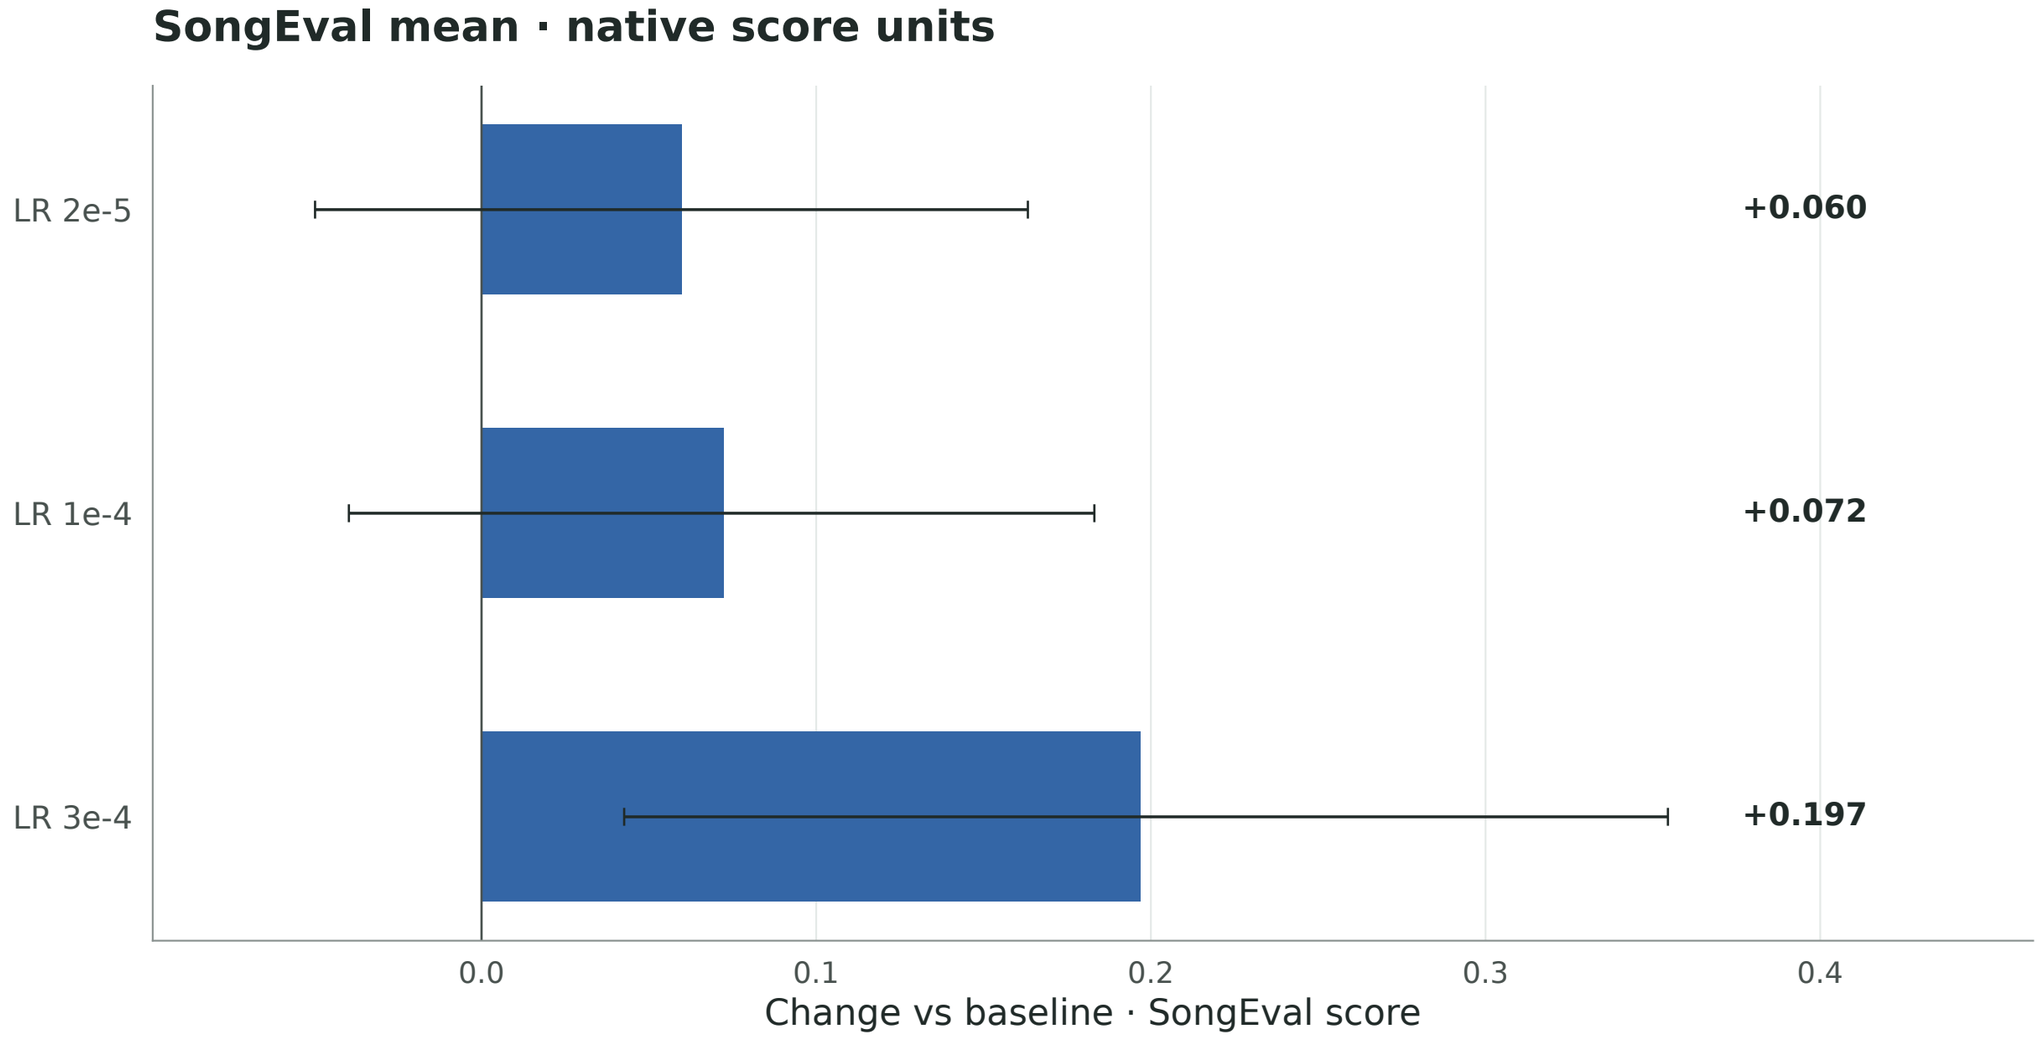

N=30 paired conditions · 95% condition-bootstrap CI · 10,000 resamples · one training seed

Beta 0.01 · rank-4 q/v LoRA · 600 semantic-token cap. LR selection and intervals are exploratory.

Validation, not full-song test or human preference. Unadjusted intervals; no across-model causal ranking.